

jupyter_eagles_final

March 24, 2025

1 Exported to PDF on March 24th, 2025

2 Original Submission Date: October, 2023

3 Cursory Analysis into the Philadelphia Eagles Football Team 2013-2022

3.1 Attempting to ascertain variables that lead to wins via Margin of Victory

3.1.1 By: Abdelrahim, Omer | Musabandesu, George | Reyhan, Achmed | O'Keefe, Connor

```
[5]: # Libraries Install
from fitter import Fitter
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import scipy as sp
import statsmodels.api as sm
import statsmodels.formula.api as smf
import pandas as pd
import scipy.stats as stats
import requests
from ydata_profiling import ProfileReport
```

```
[9]: # Cleaning Data
# df_cleaned = combined_df.dropna(axis=0, thresh=20)
# df_cleaned['Offense']=df_cleaned['Offense'].fillna(0)
# df_cleaned['Defense']=df_cleaned['Defense'].fillna(0)
# df_cleaned['Win']=np.where(df_cleaned['Unnamed: 5_level_0']=='W',1,0)
# df_cleaned['Unnamed: 8_level_0']=df_cleaned['Unnamed: 8_level_0'].fillna(0)
# df_cleaned['Home_Games']=np.where(df_cleaned['Unnamed: 8_level_0']=='@',1,0)
# df_cleaned=df_cleaned.drop(columns=['Unnamed: 8_level_0', 'Unnamed: 6_level_0'], axis=1)
```

```
[10]: # df_cleaned.to_csv('Eagles_Data_Final.csv')
# Some header cleaning such that ther are no subheaders
```

```
# eagles=pd.read_csv("/Users/omerabdelrahim/Downloads/Eagles_Data_Final.csv",
↳sep=";")
# eagles['MOV']=eagles['TmScore']-eagles['OppScore']
# eagles['MOT']=eagles['DTO']-eagles['OTO']
# eagles=eagles.drop(['Rec'], axis=1)
# eagles.to_csv('Eagles_10yr.csv')
```

4 1. Descriptive Analysis: Perform a univariate analysis following the steps below.

We will be analyzing data for the Philadelphia Eagles Football Team from the years 2013 to 2022. Our goal is to ultimately find out the most appropriate variables that contributed to Margin Of Victory (MOV) for the team.

4.1 (a) Begin by providing a descriptive analysis of your variables (include all predictors and response variable). This should include things like histograms, quantile plots, correlation plots, etc.

4.1.1 Variable Names:

MOV= Margin of Victory

MOT = Margin of Turnover

O1stD = First Downs Gained

D1stD = First Downs Allowed

DRushY = Rush Yards Allowed

ORushY = Rush Yards Gained

Home_Games = Indicator Variable for whether the game was played at home. 0 means away, 1 means home

EOffense = Expected Offensive Points

EDefense = Expected Defensive Points

ESpTms = Expected Specials Teams Points

```
[6]: eagles=pd.read_csv('Eagles_10yr (1).csv')
eagles.describe()
```

```
[6]:      Unnamed: 0  Observations      Year      TmScore      OppScore  \
count  173.000000    173.000000    173.000000  173.000000  173.000000
mean    86.000000    86.000000    2017.624277   25.433526   22.531792
std    50.084928    50.084928     2.900058    9.385319   10.502789
min      0.000000     0.000000    2013.000000     0.000000     0.000000
25%    43.000000    43.000000    2015.000000    20.000000    16.000000
50%    86.000000    86.000000    2018.000000    24.000000    22.000000
75%   129.000000   129.000000    2020.000000    32.000000    27.000000
```

max	172.000000	172.000000	2022.000000	54.000000	53.000000	
-----	------------	------------	-------------	-----------	-----------	--

	O1stD	OTotYd	OPassY	ORushY	OTO	D1stD \
count	173.000000	173.000000	173.000000	173.000000	173.000000	173.000000
mean	21.566474	368.005780	239.242775	128.763006	1.427746	20.144509
std	4.403015	82.135983	75.380423	58.052567	1.249062	5.293363
min	9.000000	139.000000	82.000000	7.000000	0.000000	8.000000
25%	19.000000	313.000000	179.000000	94.000000	0.000000	16.000000
50%	22.000000	373.000000	232.000000	118.000000	1.000000	20.000000
75%	25.000000	428.000000	298.000000	156.000000	2.000000	24.000000
max	34.000000	542.000000	462.000000	363.000000	5.000000	35.000000

	DTotYd	DPassY	DRushY	DTO	Win	Home_Games \
count	173.000000	173.000000	173.000000	173.000000	173.000000	173.000000
mean	350.965318	243.445087	107.520231	1.456647	0.566474	0.485549
std	93.325630	82.640041	46.532758	1.183523	0.497000	0.501242
min	89.000000	61.000000	6.000000	0.000000	0.000000	0.000000
25%	287.000000	182.000000	74.000000	1.000000	0.000000	0.000000
50%	349.000000	239.000000	101.000000	1.000000	1.000000	0.000000
75%	412.000000	297.000000	134.000000	2.000000	1.000000	1.000000
max	613.000000	500.000000	283.000000	5.000000	1.000000	1.000000

	MOV	MOT
count	173.000000	173.000000
mean	2.901734	0.028902
std	14.700366	1.789594
min	-41.000000	-5.000000
25%	-6.000000	-1.000000
50%	3.000000	0.000000
75%	11.000000	1.000000
max	43.000000	5.000000

```
[7]: # Replace the commas in the data with periods
eagles['EOffense'] = eagles.EOffense.str.replace(',', '.')
eagles['EDefense'] = eagles.EDefense.str.replace(',', '.')
eagles['ESpTms'] = eagles.ESpTms.str.replace(',', '.')
```

```
[8]: # Convert the objects to floats
eagles['EOffense'] = eagles.EOffense.astype(float)
eagles['EDefense'] = eagles.EDefense.astype(float)
eagles['ESpTms'] = eagles.ESpTms.astype(float)
```

```
[9]: print(eagles.columns)
```

```
Index(['Unnamed: 0', 'Observations', 'Year', 'Week', 'Day', 'Date', 'Time',
      'Source', 'WL', 'Opp', 'TmScore', 'OppScore', 'O1stD', 'OTotYd',
      'OPassY', 'ORushY', 'OTO', 'D1stD', 'DTotYd', 'DPassY', 'DRushY', 'DTO',
      'EOffense', 'EDefense', 'ESpTms', 'Win', 'Home_Games', 'MOV', 'MOT'],
```

```
dtype='object')
```

```
[10]: print(type(eagles))
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
[11]: # Check for null values
eagles.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 173 entries, 0 to 172
```

```
Data columns (total 29 columns):
```

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	173 non-null	int64
1	Observations	173 non-null	int64
2	Year	173 non-null	int64
3	Week	173 non-null	object
4	Day	173 non-null	object
5	Date	173 non-null	object
6	Time	173 non-null	object
7	Source	173 non-null	object
8	WL	173 non-null	object
9	Opp	173 non-null	object
10	TmScore	173 non-null	int64
11	OppScore	173 non-null	int64
12	O1stD	173 non-null	int64
13	OTotYd	173 non-null	int64
14	OPassY	173 non-null	int64
15	ORushY	173 non-null	int64
16	OTO	173 non-null	int64
17	D1stD	173 non-null	int64
18	DTotYd	173 non-null	int64
19	DPassY	173 non-null	int64
20	DRushY	173 non-null	int64
21	DTO	173 non-null	int64
22	EOffense	173 non-null	float64
23	EDefense	173 non-null	float64
24	ESpTms	173 non-null	float64
25	Win	173 non-null	int64
26	Home_Games	173 non-null	int64
27	MOV	173 non-null	int64
28	MOT	173 non-null	int64

```
dtypes: float64(3), int64(19), object(7)
```

```
memory usage: 39.3+ KB
```

```
[12]: # Count the number of null values
eagles.isnull().sum()
```

```
[12]: Unnamed: 0      0
      Observations    0
      Year            0
      Week            0
      Day             0
      Date            0
      Time            0
      Source          0
      WL              0
      Opp             0
      TmScore         0
      OppScore        0
      O1stD           0
      OTotYd          0
      OPassY          0
      ORushY          0
      OT0             0
      D1stD           0
      DTotYd          0
      DPassY          0
      DRushY          0
      DT0             0
      EOffense        0
      EDefense        0
      ESpTms          0
      Win             0
      Home_Games      0
      MOV             0
      MOT             0
      dtype: int64
```

```
[14]: # Initial linear model test
model1 = smf.ols(formula='MOV~MOT+DRushY+ORushY+O1stD+D1stD+Home_Games',
↳data=eagles)
model1_results=model1.fit()
print(model1_results.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          MOV      R-squared:                0.600
Model:                  OLS      Adj. R-squared:           0.585
Method:                 Least Squares      F-statistic:           41.45
Date:                   Tue, 24 Oct 2023     Prob (F-statistic):     1.30e-30
Time:                   23:49:59      Log-Likelihood:         -630.78
No. Observations:      173      AIC:                   1276.
Df Residuals:          166      BIC:                   1298.
Df Model:               6
Covariance Type:       nonrobust
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	15.0746	5.268	2.861	0.005	4.673	25.476
MOT	3.8123	0.417	9.151	0.000	2.990	4.635
DRushY	-0.0360	0.018	-2.022	0.045	-0.071	-0.001
ORushY	0.0441	0.014	3.058	0.003	0.016	0.073
O1stD	0.3241	0.189	1.715	0.088	-0.049	0.697
D1stD	-1.0093	0.159	-6.366	0.000	-1.322	-0.696
Home_Games	-1.5640	1.478	-1.058	0.291	-4.481	1.353
Omnibus:		3.602	Durbin-Watson:			1.869
Prob(Omnibus):		0.165	Jarque-Bera (JB):			3.141
Skew:		0.282	Prob(JB):			0.208
Kurtosis:		3.343	Cond. No.			1.30e+03

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.3e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Our initial regression suggests multicollinearity (there are many endogenous variables). The offensive variables all have positive correlation with MOV, and the defensive vice a versa. The intercept starts at 15 MOV, which seems rather ridiculous and is probably not interpretable. All p-values are significant except home games which has a bizarre negative value. Maybe a logit regression or a 2 stage least squares would be appropriate for such an identifier variable rather than normal regression.

```
[15]: # Histograms and Density plots

plt.figure(figsize = (10,6))
sns.histplot(eagles.MOV, stat = "density")
sns.kdeplot(eagles.MOV, color = "red")
plt.title("Philadelphia Eagles Margin of Victory 2013-2022")
plt.xlabel("Margin of Victory")
plt.show()

plt.figure(figsize = (10,6))
sns.histplot(eagles.MOT, stat = "density")
sns.kdeplot(eagles.MOT, color = "red")
plt.title("Philadelphia Eagles Margin of Turnovers 2013-2022")
plt.xlabel("Margin of Turnovers")
plt.show()

plt.figure(figsize = (10,6))
sns.histplot(eagles.O1stD, stat = "density")
```

```


sns.kdeplot(eagles.O1stD, color = "red")
plt.title("Philadelphia Eagles 1st Downs Gained 2013-2022")
plt.xlabel("1st Downs Gained")
plt.show()



plt.figure(figsize = (10,6))
sns.histplot(eagles.D1stD, stat = "density")
sns.kdeplot(eagles.D1stD, color = "red")
plt.title("Philadelphia Eagles 1st Downs Allowed 2013-2022")
plt.xlabel("1st Downs Allowed")
plt.show()



plt.figure(figsize = (10,6))
sns.histplot(eagles.DRushY, stat = "density")
sns.kdeplot(eagles.DRushY, color = "red")
plt.title("Philadelphia Eagles Rush Yards Allowed 2013-2022")
plt.xlabel("Rush Yards Allowed")
plt.show()


plt.figure(figsize = (10,6))
sns.histplot(eagles.ORushY, stat = "density")
sns.kdeplot(eagles.ORushY, color = "red")
plt.title("Philadelphia Eagles Rush Yards Gained 2013-2022")
plt.xlabel("Rush Yards Gained")
plt.show()

```









The density curve on the above plots appears to show a normal distribution. The variables DRushY and ORushY appear to show more skewness compared to the rest of the variables. These might need to be transformed.

```
[16]: # QQ-Plots

stats.probplot(eagles.MOV, dist="norm", plot=plt)
plt.title("Philadelphia Eagles Margin of Victory 2013-2022")
plt.show()

stats.probplot(eagles.MOT, dist="norm", plot=plt)
plt.title("Philadelphia Eagles Margin of Turnover 2013-2022")
plt.show()


stats.probplot(eagles.O1stD, dist="norm", plot=plt)
plt.title("Philadelphia Eagles Margin of Turnover 2013-2022")
plt.show()


stats.probplot(eagles.D1stD, dist="norm", plot=plt)
plt.title("Philadelphia Eagles Margin of Turnover 2013-2022")
plt.show()


stats.probplot(eagles.DRushY, dist="norm", plot=plt)
plt.title("Philadelphia Eagles Margin of Turnover 2013-2022")
```


```
plt.show()


stats.probplot(eagles.ORushY, dist="norm", plot=plt)
plt.title("Philadelphia Eagles Margin of Turnover 2013-2022")
plt.show()
```














The data in the six variables appears to be normal for the most part even though some deviations from the $x=y$ line can be observed. More quantitative analysis is needed to conclude normality.

```
[17]: # Correlation Plot
r_vars=eagles[['MOT', 'MOV', 'O1stD', 'D1stD', 'ORushY', 'DRushY', 'Home_Games']]

plt.figure(figsize=(13,7))
data=r_vars
c= data.corr()
sns.heatmap(c, cmap="Blues", annot=True, square = True)
```


```
[17]: <Axes: >
```




As can be viewed in this correlation plot “heatmap” some of the strongest relative correlations are between MOV and MOT as well as between O1stD and ORushY and vice a versa for the defensive statistics

```
[14]: # Pair Plot
sns.pairplot(r_vars, kind='reg')
```

```
[14]: <seaborn.axisgrid.PairGrid at 0x233c6d394d0>
```



4.2 (b) Discuss your findings from doing an exploratory analysis using Pandas Profiling. Did you discover anything new?

```
[15]: ProfileReport(r_vars)
```

```
Summarize dataset:  0%|          | 0/5 [00:00<?, ?it/s]
```

```
Generate report structure:  0%|          | 0/1 [00:00<?, ?it/s]
```

```
Render HTML:  0%|          | 0/1 [00:00<?, ?it/s]
```

```
<IPython.core.display.HTML object>
```

```
[15]:
```

The report shows results similar to what we found in 1(a) e.g., number of variables, number of observations and missing cells etc. One new discovery is that the report states that the variable MOT is highly overall correlated with MOV and vice versa. We noticed from our correlation plot that the value is 0.55

4.3 (c) Estimate density distributions (e.g., Cullen & Frey) for all your variables, and show the plots with the respective fits.

```
[26]: from fitter import Fitter
      f = Fitter(eagles.MOV)
      f.fit()
      f.summary()
```

```
SKIPPED _fit distribution (taking more than 30 seconds)
SKIPPED kstwo distribution (taking more than 30 seconds)
SKIPPED loguniform distribution (taking more than 30 seconds)
SKIPPED rv_histogram distribution (taking more than 30 seconds)
SKIPPED rv_continuous distribution (taking more than 30 seconds)
SKIPPED reciprocal distribution (taking more than 30 seconds)
SKIPPED genexpon distribution (taking more than 30 seconds)
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
IntegrationWarning: The maximum number of subdivisions (50) has been achieved.
    If increasing the limit yields no improvement it is advised to analyze
    the integrand in order to determine the difficulties.  If the position of a
    local difficulty can be determined (singularity, discontinuity) one will
    probably gain from splitting up the interval and calling the integrator
    on the subranges.  Perhaps a special-purpose integrator should be used.
    quad_r = quad(f, low, high, args=args, full_output=self.full_output,
SKIPPED kappa4 distribution (taking more than 30 seconds)
SKIPPED levy_stable distribution (taking more than 30 seconds)
SKIPPED ncx2 distribution (taking more than 30 seconds)
SKIPPED recipinvgauss distribution (taking more than 30 seconds)
SKIPPED studentized_range distribution (taking more than 30 seconds)
SKIPPED vonmises distribution (taking more than 30 seconds)
SKIPPED vonmises_line distribution (taking more than 30 seconds)
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
IntegrationWarning: The integral is probably divergent, or slowly convergent.
    quad_r = quad(f, low, high, args=args, full_output=self.full_output,
```

```
[26]:
```


	sumsquare_error	aic	bic	kl_div	ks_statistic \
dweibull	0.016563	1002.625252	1012.085127	inf	0.065805
genlogistic	0.016883	1003.542444	1013.002318	inf	0.050656
mielke	0.016904	1000.848890	1013.462057	inf	0.050348
gengamma	0.016909	1049.243331	1061.856497	inf	0.064808
burr	0.016924	1001.798222	1014.411388	inf	0.050933

	ks_pvalue
dweibull	0.423897

```

genlogistic 0.746749
mielke       0.753335
gengamma     0.443118
burr         0.740804

```



```

[27]: f = Fitter(eagles.MOT)
      f.fit()
      f.summary()

```

```

SKIPPED _fit distribution (taking more than 30 seconds)
SKIPPED kstwo distribution (taking more than 30 seconds)
SKIPPED loguniform distribution (taking more than 30 seconds)
SKIPPED rv_continuous distribution (taking more than 30 seconds)
SKIPPED rv_histogram distribution (taking more than 30 seconds)
SKIPPED reciprocal distribution (taking more than 30 seconds)
SKIPPED genexpon distribution (taking more than 30 seconds)
SKIPPED kappa4 distribution (taking more than 30 seconds)
SKIPPED levy_stable distribution (taking more than 30 seconds)
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
IntegrationWarning: The maximum number of subdivisions (50) has been achieved.
    If increasing the limit yields no improvement it is advised to analyze
    the integrand in order to determine the difficulties.  If the position of a

```


local difficulty can be determined (singularity, discontinuity) one will probably gain from splitting up the interval and calling the integrator on the subranges. Perhaps a special-purpose integrator should be used.

```
quad_r = quad(f, low, high, args=args, full_output=self.full_output,
SKIPPED ncx2 distribution (taking more than 30 seconds)
SKIPPED recipinvgauss distribution (taking more than 30 seconds)
SKIPPED studentized_range distribution (taking more than 30 seconds)
SKIPPED vonmises distribution (taking more than 30 seconds)
SKIPPED vonmises_line distribution (taking more than 30 seconds)
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
IntegrationWarning: The integral is probably divergent, or slowly convergent.
quad_r = quad(f, low, high, args=args, full_output=self.full_output,
```

```
[27]:
```

	sumsquare_error	aic	bic	kl_div	ks_statistic \
dgamma	13.796826	633.603003	643.062878	inf	0.245257
dweibull	13.971118	573.850875	583.310750	inf	0.175129
gennorm	14.141304	569.479033	578.938908	inf	0.110779
weibull_min	14.141599	566.108074	575.567949	inf	0.107337
gengamma	14.142542	568.959919	581.573085	inf	0.111033

	ks_pvalue
dgamma	1.194577e-09
dweibull	4.123580e-05
gennorm	2.643214e-02
weibull_min	3.438646e-02
gengamma	2.591388e-02



```
[28]: f = Fitter(eagles.E0ffense)
      f.fit()
      f.summary()
```

```
SKIPPED _fit distribution (taking more than 30 seconds)
SKIPPED kstwo distribution (taking more than 30 seconds)
SKIPPED loguniform distribution (taking more than 30 seconds)
SKIPPED rv_continuous distribution (taking more than 30 seconds)
SKIPPED rv_histogram distribution (taking more than 30 seconds)
SKIPPED reciprocal distribution (taking more than 30 seconds)
SKIPPED genexpon distribution (taking more than 30 seconds)
SKIPPED kappa4 distribution (taking more than 30 seconds)
SKIPPED levy_stable distribution (taking more than 30 seconds)
SKIPPED ncx2 distribution (taking more than 30 seconds)
SKIPPED recipinvgauss distribution (taking more than 30 seconds)
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
IntegrationWarning: The maximum number of subdivisions (50) has been achieved.
    If increasing the limit yields no improvement it is advised to analyze
    the integrand in order to determine the difficulties.  If the position of a
    local difficulty can be determined (singularity, discontinuity) one will
    probably gain from splitting up the interval and calling the integrator
    on the subranges.  Perhaps a special-purpose integrator should be used.
```


```
quad_r = quad(f, low, high, args=args, full_output=self.full_output,
SKIPPED t distribution (taking more than 30 seconds)
SKIPPED studentized_range distribution (taking more than 30 seconds)
SKIPPED vonmises distribution (taking more than 30 seconds)
SKIPPED vonmises_line distribution (taking more than 30 seconds)
```

```
[28]:
```

	sumsquare_error	aic	bic	kl_div	ks_statistic \
burr12	0.017307	846.461690	859.074856	inf	0.033253
genhyperbolic	0.017336	848.417290	864.183748	inf	0.026649
skewnorm	0.017338	844.712977	854.172851	inf	0.027018
pearson3	0.017338	844.221344	853.681218	inf	0.026350
powernorm	0.017341	844.674398	854.134273	inf	0.027154

	ks_pvalue
burr12	0.987656
genhyperbolic	0.999432
skewnorm	0.999281
pearson3	0.999533
powernorm	0.999218

```
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
IntegrationWarning: The integral is probably divergent, or slowly convergent.
quad_r = quad(f, low, high, args=args, full_output=self.full_output,
```



```
[29]: f = Fitter(eagles.EDefense)
      f.fit()
      f.summary()
```


```
SKIPPED _fit distribution (taking more than 30 seconds)
SKIPPED kstwo distribution (taking more than 30 seconds)
SKIPPED loguniform distribution (taking more than 30 seconds)
SKIPPED rv_continuous distribution (taking more than 30 seconds)
SKIPPED rv_histogram distribution (taking more than 30 seconds)
SKIPPED burr distribution (taking more than 30 seconds)
SKIPPED reciprocal distribution (taking more than 30 seconds)
SKIPPED genexpon distribution (taking more than 30 seconds)
SKIPPED kappa4 distribution (taking more than 30 seconds)
SKIPPED levy_stable distribution (taking more than 30 seconds)
SKIPPED ncx2 distribution (taking more than 30 seconds)
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
IntegrationWarning: The maximum number of subdivisions (50) has been achieved.
    If increasing the limit yields no improvement it is advised to analyze
    the integrand in order to determine the difficulties.  If the position of a
    local difficulty can be determined (singularity, discontinuity) one will
    probably gain from splitting up the interval and calling the integrator
    on the subranges.  Perhaps a special-purpose integrator should be used.
    quad_r = quad(f, low, high, args=args, full_output=self.full_output,
SKIPPED recipinvgauss distribution (taking more than 30 seconds)
SKIPPED studentized_range distribution (taking more than 30 seconds)
SKIPPED t distribution (taking more than 30 seconds)
SKIPPED triang distribution (taking more than 30 seconds)
SKIPPED vonmises distribution (taking more than 30 seconds)
SKIPPED vonmises_line distribution (taking more than 30 seconds)
```

```
[29]:
```

	sumsquare_error	aic	bic	kl_div	ks_statistic \
burr12	0.009213	954.313035	966.926202	inf	0.036434
skewnorm	0.009227	955.699446	965.159321	inf	0.034841
genhyperbolic	0.009230	959.717410	975.483868	inf	0.035541
powernorm	0.009233	955.390752	964.850626	inf	0.034770
norminvgauss	0.009239	957.311268	969.924435	inf	0.035453

	ks_pvalue
burr12	0.969344
skewnorm	0.979919
genhyperbolic	0.975640
powernorm	0.980322
norminvgauss	0.976210

```
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
IntegrationWarning: The integral is probably divergent, or slowly convergent.
    quad_r = quad(f, low, high, args=args, full_output=self.full_output,
```

```
[30]: f = Fitter(eagles.ESpTms)
      f.fit()
      f.summary()
```

```
SKIPPED _fit distribution (taking more than 30 seconds)
SKIPPED kstwo distribution (taking more than 30 seconds)
SKIPPED loguniform distribution (taking more than 30 seconds)
SKIPPED crystalball distribution (taking more than 30 seconds)
SKIPPED fisk distribution (taking more than 30 seconds)
SKIPPED genexpon distribution (taking more than 30 seconds)
SKIPPED rv_continuous distribution (taking more than 30 seconds)
SKIPPED rv_histogram distribution (taking more than 30 seconds)
SKIPPED reciprocal distribution (taking more than 30 seconds)
SKIPPED kappa4 distribution (taking more than 30 seconds)
SKIPPED levy_stable distribution (taking more than 30 seconds)
SKIPPED ncx2 distribution (taking more than 30 seconds)
SKIPPED recipinvgauss distribution (taking more than 30 seconds)
SKIPPED studentized_range distribution (taking more than 30 seconds)
SKIPPED triang distribution (taking more than 30 seconds)
SKIPPED vonmises distribution (taking more than 30 seconds)
SKIPPED vonmises_line distribution (taking more than 30 seconds)
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
```

IntegrationWarning: The maximum number of subdivisions (50) has been achieved.

If increasing the limit yields no improvement it is advised to analyze the integrand in order to determine the difficulties. If the position of a local difficulty can be determined (singularity, discontinuity) one will probably gain from splitting up the interval and calling the integrator on the subranges. Perhaps a special-purpose integrator should be used.

```
quad_r = quad(f, low, high, args=args, full_output=self.full_output,
```

```
[30]:
```


	sumsquare_error	aic	bic	kl_div	ks_statistic \
gennorm	0.079485	773.141991	782.601866	inf	0.051626
logistic	0.080359	774.586283	780.892866	inf	0.044246
tukeylambda	0.080565	776.274220	785.734095	inf	0.046655
hypsecant	0.080598	772.908962	779.215545	inf	0.054634
t	0.080683	777.893214	787.353089	inf	0.047434

	ks_pvalue
gennorm	0.725811
logistic	0.872087
tukeylambda	0.828463
hypsecant	0.659486
t	0.813315

C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate_quadpack_py.py:1225:

IntegrationWarning: The integral is probably divergent, or slowly convergent.

```
quad_r = quad(f, low, high, args=args, full_output=self.full_output,
```



```
[31]: f = Fitter(eagles.DRushY)
      f.fit()
      f.summary()
```

```
SKIPPED _fit distribution (taking more than 30 seconds)
SKIPPED kstwo distribution (taking more than 30 seconds)
SKIPPED burr distribution (taking more than 30 seconds)
SKIPPED crystalball distribution (taking more than 30 seconds)
SKIPPED fisk distribution (taking more than 30 seconds)
SKIPPED genexpon distribution (taking more than 30 seconds)
SKIPPED rv_histogram distribution (taking more than 30 seconds)
SKIPPED rv_continuous distribution (taking more than 30 seconds)
SKIPPED kappa4 distribution (taking more than 30 seconds)
SKIPPED ncx2 distribution (taking more than 30 seconds)
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
IntegrationWarning: The maximum number of subdivisions (50) has been achieved.
    If increasing the limit yields no improvement it is advised to analyze
    the integrand in order to determine the difficulties.  If the position of a
    local difficulty can be determined (singularity, discontinuity) one will
    probably gain from splitting up the interval and calling the integrator
    on the subranges.  Perhaps a special-purpose integrator should be used.
    quad_r = quad(f, low, high, args=args, full_output=self.full_output,
SKIPPED recipinvgauss distribution (taking more than 30 seconds)
SKIPPED studentized_range distribution (taking more than 30 seconds)
SKIPPED vonmises_line distribution (taking more than 30 seconds)
SKIPPED vonmises distribution (taking more than 30 seconds)
```

```
[31]:
```

	sumsquare_error	aic	bic	kl_div	\
genhyperbolic	0.000606	1273.315331	1289.081789	inf	
invweibull	0.000616	1262.841627	1272.301501	inf	
gumbel_r	0.000616	1260.772437	1267.079021	inf	
skewnorm	0.000617	1274.563632	1284.023506	inf	
norminvgauss	0.000621	1271.314025	1283.927192	inf	

	ks_statistic	ks_pvalue
genhyperbolic	0.035634	0.975024
invweibull	0.032937	0.988895
gumbel_r	0.032960	0.988809
skewnorm	0.034656	0.980961
norminvgauss	0.038743	0.948471


```
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
IntegrationWarning: The integral is probably divergent, or slowly convergent.
    quad_r = quad(f, low, high, args=args, full_output=self.full_output,
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
IntegrationWarning: The occurrence of roundoff error is detected, which prevents
```

the requested tolerance from being achieved. The error may be underestimated.

```
quad_r = quad(f, low, high, args=args, full_output=self.full_output,  
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:  
IntegrationWarning: The maximum number of subdivisions (50) has been achieved.
```

If increasing the limit yields no improvement it is advised to analyze the integrand in order to determine the difficulties. If the position of a local difficulty can be determined (singularity, discontinuity) one will probably gain from splitting up the interval and calling the integrator on the subranges. Perhaps a special-purpose integrator should be used.

```
quad_r = quad(f, low, high, args=args, full_output=self.full_output,
```



```
[32]: f = Fitter(eagles.ORushY)  
f.fit()  
f.summary()
```

```
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:  
IntegrationWarning: The occurrence of roundoff error is detected, which prevents  
the requested tolerance from being achieved. The error may be  
underestimated.
```

```
quad_r = quad(f, low, high, args=args, full_output=self.full_output,  
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:  
IntegrationWarning: The maximum number of subdivisions (50) has been achieved.
```

If increasing the limit yields no improvement it is advised to analyze the integrand in order to determine the difficulties. If the position of a local difficulty can be determined (singularity, discontinuity) one will probably gain from splitting up the interval and calling the integrator on the subranges. Perhaps a special-purpose integrator should be used.

```
quad_r = quad(f, low, high, args=args, full_output=self.full_output,
```

SKIPPED _fit distribution (taking more than 30 seconds)

SKIPPED burr distribution (taking more than 30 seconds)

SKIPPED crystalball distribution (taking more than 30 seconds)

SKIPPED kstwo distribution (taking more than 30 seconds)

SKIPPED fisk distribution (taking more than 30 seconds)

SKIPPED genexpon distribution (taking more than 30 seconds)

SKIPPED rv_continuous distribution (taking more than 30 seconds)

SKIPPED rv_histogram distribution (taking more than 30 seconds)

SKIPPED kappa4 distribution (taking more than 30 seconds)

SKIPPED levy_stable distribution (taking more than 30 seconds)

SKIPPED ncx2 distribution (taking more than 30 seconds)

SKIPPED recipinvgauss distribution (taking more than 30 seconds)

SKIPPED t distribution (taking more than 30 seconds)

SKIPPED studentized_range distribution (taking more than 30 seconds)

SKIPPED triang distribution (taking more than 30 seconds)

SKIPPED vonmises distribution (taking more than 30 seconds)

SKIPPED vonmises_line distribution (taking more than 30 seconds)

C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate_quadpack_py.py:1225:
IntegrationWarning: The maximum number of subdivisions (50) has been achieved.

If increasing the limit yields no improvement it is advised to analyze the integrand in order to determine the difficulties. If the position of a local difficulty can be determined (singularity, discontinuity) one will probably gain from splitting up the interval and calling the integrator on the subranges. Perhaps a special-purpose integrator should be used.

```
quad_r = quad(f, low, high, args=args, full_output=self.full_output,
```

C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate_quadpack_py.py:1225:
IntegrationWarning: The occurrence of roundoff error is detected, which prevents the requested tolerance from being achieved. The error may be underestimated.

```
quad_r = quad(f, low, high, args=args, full_output=self.full_output,
```

C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate_quadpack_py.py:1225:
IntegrationWarning: The occurrence of roundoff error is detected, which prevents the requested tolerance from being achieved. The error may be underestimated.

```
quad_r = quad(f, low, high, args=args, full_output=self.full_output,
```

C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate_quadpack_py.py:1225:
IntegrationWarning: The maximum number of subdivisions (50) has been achieved.

If increasing the limit yields no improvement it is advised to analyze the integrand in order to determine the difficulties. If the position of a local difficulty can be determined (singularity, discontinuity) one will probably gain from splitting up the interval and calling the integrator on the subranges. Perhaps a special-purpose integrator should be used.

```
quad_r = quad(f, low, high, args=args, full_output=self.full_output,
```

```
[32]:
```


	sumsquare_error	aic	bic	kl_div	ks_statistic \
loglaplace	0.000275	1332.779068	1342.238943	inf	0.034855
dgamma	0.000285	1369.938881	1379.398755	inf	0.066331
dweibull	0.000289	1375.764973	1385.224847	inf	0.071209
laplace	0.000290	1371.085432	1377.392016	inf	0.063825
gennorm	0.000292	1375.263816	1384.723691	inf	0.062297

	ks_pvalue
loglaplace	0.979839
dgamma	0.413965
dweibull	0.328583
laplace	0.462500
gennorm	0.493461

C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate_quadpack_py.py:1225:

IntegrationWarning: The integral is probably divergent, or slowly convergent.

```
quad_r = quad(f, low, high, args=args, full_output=self.full_output,
```



```
[33]: f = Fitter(eagles.01stD)
f.fit()
f.summary()
```

```

C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
IntegrationWarning: The maximum number of subdivisions (50) has been achieved.
    If increasing the limit yields no improvement it is advised to analyze
    the integrand in order to determine the difficulties.  If the position of a
    local difficulty can be determined (singularity, discontinuity) one will
    probably gain from splitting up the interval and calling the integrator
    on the subranges.  Perhaps a special-purpose integrator should be used.
    quad_r = quad(f, low, high, args=args, full_output=self.full_output,
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
IntegrationWarning: The occurrence of roundoff error is detected, which prevents
    the requested tolerance from being achieved.  The error may be
    underestimated.
    quad_r = quad(f, low, high, args=args, full_output=self.full_output,
SKIPPED _fit distribution (taking more than 30 seconds)
SKIPPED burr distribution (taking more than 30 seconds)
SKIPPED crystalball distribution (taking more than 30 seconds)
SKIPPED kstwo distribution (taking more than 30 seconds)
SKIPPED fisk distribution (taking more than 30 seconds)
C:\Users\kabde\anaconda3\Lib\site-
packages\scipy\stats\_continuous_distns.py:3485: IntegrationWarning: The
algorithm does not converge. Roundoff error is detected
    in the extrapolation table.  It is assumed that the requested tolerance
    cannot be achieved, and that the returned result (if full_output = 1) is
    the best which can be obtained.
    t1 = integrate.quad(llc, -np.inf, x)[0]
SKIPPED genexpon distribution (taking more than 30 seconds)
SKIPPED rv_continuous distribution (taking more than 30 seconds)
SKIPPED rv_histogram distribution (taking more than 30 seconds)
SKIPPED kappa4 distribution (taking more than 30 seconds)
SKIPPED ncx2 distribution (taking more than 30 seconds)
SKIPPED recipinvgauss distribution (taking more than 30 seconds)
SKIPPED studentized_range distribution (taking more than 30 seconds)
SKIPPED t distribution (taking more than 30 seconds)
SKIPPED triang distribution (taking more than 30 seconds)
SKIPPED vonmises distribution (taking more than 30 seconds)
SKIPPED vonmises_line distribution (taking more than 30 seconds)
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
IntegrationWarning: The maximum number of subdivisions (50) has been achieved.
    If increasing the limit yields no improvement it is advised to analyze
    the integrand in order to determine the difficulties.  If the position of a
    local difficulty can be determined (singularity, discontinuity) one will
    probably gain from splitting up the interval and calling the integrator
    on the subranges.  Perhaps a special-purpose integrator should be used.
    quad_r = quad(f, low, high, args=args, full_output=self.full_output,
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
IntegrationWarning: The occurrence of roundoff error is detected, which prevents
    the requested tolerance from being achieved.  The error may be
    underestimated.

```

```
quad_r = quad(f, low, high, args=args, full_output=self.full_output,
```

```
[33]:
```

	sumsquare_error	aic	bic	kl_div	ks_statistic \
skewnorm	0.827303	768.117380	777.577254	inf	0.056435
genhyperbolic	0.827397	772.006679	787.773137	inf	0.101654
weibull_min	0.827702	766.722469	776.182343	inf	0.064663
burr12	0.827785	767.858047	780.471213	inf	0.062137
norminvgauss	0.827786	768.963171	781.576337	inf	0.058796

	ks_pvalue
skewnorm	0.619514
genhyperbolic	0.052138
weibull_min	0.445931
burr12	0.496765
norminvgauss	0.567677

```
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
```

```
IntegrationWarning: The maximum number of subdivisions (50) has been achieved.
```

If increasing the limit yields no improvement it is advised to analyze the integrand in order to determine the difficulties. If the position of a local difficulty can be determined (singularity, discontinuity) one will probably gain from splitting up the interval and calling the integrator on the subranges. Perhaps a special-purpose integrator should be used.

```
quad_r = quad(f, low, high, args=args, full_output=self.full_output,
```

```
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
```


```
IntegrationWarning: The occurrence of roundoff error is detected, which prevents the requested tolerance from being achieved. The error may be underestimated.
```

```
quad_r = quad(f, low, high, args=args, full_output=self.full_output,
```

```
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
```

```
IntegrationWarning: The integral is probably divergent, or slowly convergent.
```

```
quad_r = quad(f, low, high, args=args, full_output=self.full_output,
```

```
[34]: f = Fitter(eagles.D1stD)
      f.fit()
      f.summary()
```

```
SKIPPED _fit distribution (taking more than 30 seconds)
SKIPPED burr distribution (taking more than 30 seconds)
SKIPPED crystalball distribution (taking more than 30 seconds)
SKIPPED kstwo distribution (taking more than 30 seconds)
SKIPPED fisk distribution (taking more than 30 seconds)
SKIPPED genexpon distribution (taking more than 30 seconds)
C:\Users\kabde\anaconda3\Lib\site-
packages\scipy\stats\_continuous_distns.py:3485: IntegrationWarning: The
integral is probably divergent, or slowly convergent.
  t1 = integrate.quad(llc, -np.inf, x)[0]
SKIPPED gengamma distribution (taking more than 30 seconds)
SKIPPED rv_continuous distribution (taking more than 30 seconds)
SKIPPED rv_histogram distribution (taking more than 30 seconds)
SKIPPED kappa4 distribution (taking more than 30 seconds)
SKIPPED johnsonsu distribution (taking more than 30 seconds)
SKIPPED levy_stable distribution (taking more than 30 seconds)
SKIPPED ncx2 distribution (taking more than 30 seconds)
SKIPPED powerlognorm distribution (taking more than 30 seconds)
```

```

SKIPPED recipinvgauss distribution (taking more than 30 seconds)
SKIPPED studentized_range distribution (taking more than 30 seconds)
SKIPPED t distribution (taking more than 30 seconds)
SKIPPED trapezoid distribution (taking more than 30 seconds)
SKIPPED triang distribution (taking more than 30 seconds)
SKIPPED vonmises_line distribution (taking more than 30 seconds)
SKIPPED vonmises distribution (taking more than 30 seconds)
C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
IntegrationWarning: The maximum number of subdivisions (50) has been achieved.

```

If increasing the limit yields no improvement it is advised to analyze the integrand in order to determine the difficulties. If the position of a local difficulty can be determined (singularity, discontinuity) one will probably gain from splitting up the interval and calling the integrator on the subranges. Perhaps a special-purpose integrator should be used.

```
quad_r = quad(f, low, high, args=args, full_output=self.full_output,
```


```
[34]:
```

	sumsquare_error	aic	bic	kl_div	ks_statistic \
dweibull	0.616115	755.600468	765.060343	inf	0.091170
dgamma	0.616229	757.357144	766.817019	inf	0.098407
genhyperbolic	0.618622	744.869611	760.636069	inf	1.000000
invgauss	0.618757	740.754785	750.214660	inf	0.067423
invgamma	0.618797	738.946339	748.406214	inf	0.062522

	ks_pvalue
dweibull	0.105878
dgamma	0.065468
genhyperbolic	0.000000
invgauss	0.393778
invgamma	0.488842

```

C:\Users\kabde\anaconda3\Lib\site-packages\scipy\integrate\_quadpack_py.py:1225:
IntegrationWarning: The integral is probably divergent, or slowly convergent.
quad_r = quad(f, low, high, args=args, full_output=self.full_output,
```



4.4 (d) Identify if there are any non-linearities within your variables. What transformations should you perform to make them linear? What would happen if you included nonlinear variables in your regression models without transforming them first?

Based on the graphical results from the pair plots as well as `panda_profiling` in section (b) the data appears to be linear. We will need to perform a transformation of the two variables `DRushY` and `ORushY`. The results of the transformation using Box-cox will be shown in the later section to follow.

```
[18]: # Box-Cox Transformation of DRushY and ORushY
import scipy

bc_DRY, lambda_DRY = scipy.stats.boxcox(eagles["DRushY"])
print(lambda_DRY)

sns.histplot(eagles["DRushY"])
plt.title("Original DRushY")
plt.show()

sns.histplot(bc_DRY)
```


```
plt.title("Box-Cox Transformed: DRushY")
plt.show()


bc_ORY,lambda_ORY = scipy.stats.boxcox(eagles["ORushY"])
print(lambda_ORY)

sns.histplot(eagles["ORushY"])
plt.title("Original ORushY")
plt.show()


sns.histplot(bc_ORY)
plt.title("Box-Cox Transformed: ORushY")
plt.show()
```


0.49829613545987844





0.4559889588939842







Upon comparison of the two variables before and after the transformation - the results show that the distributions appears normal.

```
[19]: # Here, we append the original dataset with the tranformed variables
eagles['BCORY']=bc_ORY
eagles['BCDRY']=bc_DRY
# New dataset r_vars
r_vars=eagles[['MOT', 'MOV', 'O1stD', 'D1stD', 'BCORY', 'BCDRY', 'Home_Games']]
```


```
[18]: # Graphical Represenation of the transformed variables
sns.lmplot(data=r_vars, x='BCORY', y='MOV', line_kws={'color':'red'},
           lowess=False, height=5, aspect=1)
plt.grid()
sns.lmplot(data=r_vars, x='BCDRY', y='MOV', line_kws={'color':'green'},
           lowess=False, height=5, aspect=1)
plt.grid()
```





```
[19]: # Pair Plot  
sns.pairplot(r_vars, kind='reg')
```

```
[19]: <seaborn.axisgrid.PairGrid at 0x233d61a2c90>
```



The new pair plot with the transformed variables BCDRY and BCORY shows these two variables with a normal distribution.

If non-linear variables are included in the model the results will not be robust and reliable in comparison to those of a model with linear variables.

4.5 (e) Comment on any outliers and/or unusual features of your variables.

In the graphs below we will see very few outliers (less than 5%) but we do not find it appropriate to remove these outliers due to the nature of data. We think the removal of outliers will distort the findings of our analysis.

```
[20]: # Outliers
sns.lmplot(data = eagles, x = 'MOT', y = 'MOV')
```

```
plt.title("MOT Regressed")


sns.lmplot(data = eagles, x = 'DRushY', y = 'MOV')
plt.title("DRushY Regressed")


sns.lmplot(data = eagles, x = 'ORushY', y = 'MOV')
plt.title("ORushY Regressed")

sns.lmplot(data = eagles, x = 'O1stD', y = 'MOV')
plt.title("O1stD Regressed")

sns.lmplot(data = eagles, x = 'D1stD', y = 'MOV')
plt.title("D1st Regressed")


plt.show()
```












4.6 (f) If you have any NAs, impute them using any of the methods discussed in class but make sure to justify your choice.

We do not have any NAs in our dataset as was confirmed in part 1(a)

```
[21]: # Boxplots to show outliers
plt.boxplot(eagles.MOV)
plt.grid
```

```
[21]: <function matplotlib.pyplot.grid(visible=None, which='major', axis='both',
**kwargs)>
```



The above box plot clearly indicates the presence of outliers.

5 2. Variable Selection:

5.1 (a) Using the Boruta Algorithm identify the top 2 predictors

```
[22]: from BorutaShap import BorutaShap

      from sklearn.ensemble import RandomForestRegressor


[23]: boruta_data = eagles[['MOV', 'O1stD', 'BCORY', 'D1stD', 'BCDRY', 'MOT', '
      ↪ 'Home_Games']].copy()

[24]: x = boruta_data.iloc[:, 1:]

[25]: x = boruta_data.iloc[:, 1:]
      y = boruta_data['MOV']
      Feature_Selector = BorutaShap(importance_measure='shap', classification=False)
      Feature_Selector.fit(X=x, y=y, n_trials=100, random_state=0)
      Feature_Selector.plot(which_features='all')
```

0%| | 0/100 [00:00<?, ?it/s]

5 attributes confirmed important: ['BCDRY', 'O1stD', 'D1stD', 'BCORY', 'MOT']
 1 attributes confirmed unimportant: ['Home_Games']
 0 tentative attributes remains: []



```
[26]: Feature_Selector.Subset()
```

```
[26]:
```

	BCDRY	O1stD	D1stD	BCORY	MOT
0	24.754325	27	23	23.247767	1
1	13.684295	25	20	20.182776	2
2	16.569813	21	20	13.222861	1
3	14.780713	25	13	22.923054	4
4	20.157584	24	23	18.615272	1
..	...	...	...	...	...
168	15.809353	23	11	27.313920	2
169	16.355774	25	29	13.889061	0
170	13.557671	28	18	26.859759	1
171	12.908318	21	21	18.478211	2
172	25.047405	17	26	13.981570	2

[173 rows x 5 columns]

```
[27]: from sklearn.ensemble import RandomForestClassifier
      from boruta import BorutaPy
      x = boruta_data.iloc[:, 1:]
      y = boruta_data['MOV']
      xCols = x.columns.tolist()

      currentTrainX = x.to_numpy()
      currentTrainY = y.to_numpy().ravel()

      forest = RandomForestRegressor(n_jobs=-1, max_depth = 5)
      forest.fit(currentTrainX, currentTrainY)
```

```
[27]: RandomForestRegressor(max_depth=5, n_jobs=-1)
```

```
[28]: np.int = int
      np.float = float
      np.bool = bool
      boruta = BorutaPy(forest, n_estimators='auto', verbose=2, random_state=1)
      boruta.fit(currentTrainX, currentTrainY)

      featureSupport = list((zip(xCols, boruta.support_)))
      featureSupport
```

```
Iteration:      1 / 100
Confirmed:      0
Tentative:      6
Rejected:       0
Iteration:      2 / 100
Confirmed:      0
Tentative:      6
Rejected:       0
Iteration:      3 / 100
Confirmed:      0
Tentative:      6
Rejected:       0
Iteration:      4 / 100
Confirmed:      0
Tentative:      6
Rejected:       0
Iteration:      5 / 100
Confirmed:      0
Tentative:      6
Rejected:       0
Iteration:      6 / 100
Confirmed:      0
Tentative:      6
Rejected:       0
Iteration:      7 / 100
```

```
Confirmed:      0
Tentative:      6
Rejected:       0
Iteration:      8 / 100
Confirmed:      4
Tentative:      0
Rejected:       2
```

BorutaPy finished running.

```
Iteration:      9 / 100
Confirmed:      4
Tentative:      0
Rejected:       2
```

```
[28]: [('O1stD', False),
       ('BCORY', True),
       ('D1stD', True),
       ('BCDRY', True),
       ('MOT', True),
       ('Home_Games', False)]
```

```
[29]: # Ranking Boruta results
boruta.ranking_
# for better visualization use of the boruta ranking
featureRanks = list(zip(xCols, boruta.ranking_))
sorted(featureRanks, key=lambda x: x[1])
```

```
[29]: [('BCORY', 1),
       ('D1stD', 1),
       ('BCDRY', 1),
       ('MOT', 1),
       ('O1stD', 2),
       ('Home_Games', 3)]
```

The top ranking variables from our Boruta Algorithm are the Box-Cox Transformation of Offensive (BCORY) and Defensive (BCDRY) Rushing Yards as well as MOT and Defensive 1st Downs Allowed (D1stD). The top 2 are BCORY and BCDRY

Going forward it would be prudent to work with one of these 4 explanatory variables, and we will use Mallows CP in order to come to a clearer solution.

5.2 (b) Using Mallows Cp identify the top 2 predictors

```
[31]: from RegscorePy import mallow
```

```
[21]: model = smf.ols(formula='(MOV) ~ MOT + O1stD + BCORY + D1stD + BCDRY',
    ↪data=eagles)
results = model.fit()
print(results.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  MOV    R-squared:                   0.596
Model:                        OLS    Adj. R-squared:              0.584
Method:                    Least Squares    F-statistic:                49.36
Date:                Tue, 24 Oct 2023    Prob (F-statistic):        3.40e-31
Time:                23:56:09    Log-Likelihood:            -631.49
No. Observations:                173    AIC:                        1275.
Df Residuals:                    167    BIC:                        1294.
Df Model:                        5
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	13.7689	6.117	2.251	0.026	1.692	25.846
MOT	3.8149	0.417	9.154	0.000	2.992	4.638
O1stD	0.3177	0.192	1.659	0.099	-0.060	0.696
BCORY	0.5946	0.206	2.883	0.004	0.187	1.002
D1stD	-1.0351	0.156	-6.626	0.000	-1.343	-0.727
BCDRY	-0.4035	0.186	-2.167	0.032	-0.771	-0.036

```

=====
Omnibus:                        3.207    Durbin-Watson:              1.864
Prob(Omnibus):                  0.201    Jarque-Bera (JB):            2.752
Skew:                          0.279    Prob(JB):                    0.253
Kurtosis:                      3.267    Cond. No.                    332.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Now we have no issues with multicollinearity but that's probably because of the nature of the Box-Cox transformations reducing interpretability of the variables for Rush Yards offensively and defensively. That being said we get a lower value of intercept. First downs lose their significance as well. Overall this regression doesn't really provide much valuable or interpretable data.

```
[33]: import itertools

model = smf.ols(formula='MOV ~ MOT + O1stD + BCORY + D1stD + BCDRY',
    ↪data=eagles)
results = model.fit()
y = eagles['MOV']
y_pred=results.fittedvalues
```

```

storage_cp = pd.DataFrame(columns = ["Variables", "CP"])
k = 8

for L in range(1, len(r_vars.columns[0:]) + 1):
    for subset in itertools.combinations(r_vars.columns[0:], L):

        formula1 = 'MOV~'+'+'.join(subset)

        results = smf.ols(formula=formula1, data = eagles).fit()
        y_sub = results.fittedvalues
        p = len(subset)+1

        cp = mallow.mallow(y, y_pred,y_sub, k, p)

        storage_cp = storage_cp._append({'Variables': subset, 'CP': cp},
        ignore_index = True)

```

```

[34]: pd.set_option('display.max_columns', 2)
pd.set_option('display.max_rows', 127)
storage_cp.sort_values(by = "CP")

```

```

[34]:

```

	Variables	CP
1	(MOV,)	-169.000000
17	(MOV, Home_Games)	-167.000000
15	(MOV, BCORY)	-167.000000
14	(MOV, D1stD)	-167.000000
13	(MOV, O1stD)	-167.000000
7	(MOT, MOV)	-167.000000
16	(MOV, BCDRY)	-167.000000
48	(MOV, D1stD, BCDRY)	-165.000000
47	(MOV, D1stD, BCORY)	-165.000000
46	(MOV, O1stD, Home_Games)	-165.000000
44	(MOV, O1stD, BCORY)	-165.000000
43	(MOV, O1stD, D1stD)	-165.000000
28	(MOT, MOV, O1stD)	-165.000000
29	(MOT, MOV, D1stD)	-165.000000
30	(MOT, MOV, BCORY)	-165.000000
32	(MOT, MOV, Home_Games)	-165.000000
45	(MOV, O1stD, BCDRY)	-165.000000
49	(MOV, D1stD, Home_Games)	-165.000000
31	(MOT, MOV, BCDRY)	-165.000000
51	(MOV, BCORY, Home_Games)	-165.000000
52	(MOV, BCDRY, Home_Games)	-165.000000
50	(MOV, BCORY, BCDRY)	-165.000000
87	(MOV, O1stD, BCORY, Home_Games)	-163.000000

86	(MOV, 01stD, BCORY, BCDRY)	-163.000000
65	(MOT, MOV, 01stD, BCDRY)	-163.000000
85	(MOV, 01stD, D1stD, Home_Games)	-163.000000
66	(MOT, MOV, 01stD, Home_Games)	-163.000000
67	(MOT, MOV, D1stD, BCORY)	-163.000000
84	(MOV, 01stD, D1stD, BCDRY)	-163.000000
83	(MOV, 01stD, D1stD, BCORY)	-163.000000
69	(MOT, MOV, D1stD, Home_Games)	-163.000000
70	(MOT, MOV, BCORY, BCDRY)	-163.000000
71	(MOT, MOV, BCORY, Home_Games)	-163.000000
72	(MOT, MOV, BCDRY, Home_Games)	-163.000000
68	(MOT, MOV, D1stD, BCDRY)	-163.000000
88	(MOV, 01stD, BCDRY, Home_Games)	-163.000000
63	(MOT, MOV, 01stD, D1stD)	-163.000000
90	(MOV, D1stD, BCORY, Home_Games)	-163.000000
91	(MOV, D1stD, BCDRY, Home_Games)	-163.000000
92	(MOV, BCORY, BCDRY, Home_Games)	-163.000000
64	(MOT, MOV, 01stD, BCORY)	-163.000000
89	(MOV, D1stD, BCORY, BCDRY)	-163.000000
117	(MOV, D1stD, BCORY, BCDRY, Home_Games)	-161.000000
116	(MOV, 01stD, BCORY, BCDRY, Home_Games)	-161.000000
115	(MOV, 01stD, D1stD, BCDRY, Home_Games)	-161.000000
114	(MOV, 01stD, D1stD, BCORY, Home_Games)	-161.000000
113	(MOV, 01stD, D1stD, BCORY, BCDRY)	-161.000000
107	(MOT, MOV, BCORY, BCDRY, Home_Games)	-161.000000
106	(MOT, MOV, D1stD, BCDRY, Home_Games)	-161.000000
104	(MOT, MOV, D1stD, BCORY, BCDRY)	-161.000000
103	(MOT, MOV, 01stD, BCDRY, Home_Games)	-161.000000
102	(MOT, MOV, 01stD, BCORY, Home_Games)	-161.000000
101	(MOT, MOV, 01stD, BCORY, BCDRY)	-161.000000
100	(MOT, MOV, 01stD, D1stD, Home_Games)	-161.000000
99	(MOT, MOV, 01stD, D1stD, BCDRY)	-161.000000
98	(MOT, MOV, 01stD, D1stD, BCORY)	-161.000000
105	(MOT, MOV, D1stD, BCORY, Home_Games)	-161.000000
119	(MOT, MOV, 01stD, D1stD, BCORY, BCDRY)	-159.000000
120	(MOT, MOV, 01stD, D1stD, BCORY, Home_Games)	-159.000000
121	(MOT, MOV, 01stD, D1stD, BCDRY, Home_Games)	-159.000000
122	(MOT, MOV, 01stD, BCORY, BCDRY, Home_Games)	-159.000000
123	(MOT, MOV, D1stD, BCORY, BCDRY, Home_Games)	-159.000000
125	(MOV, 01stD, D1stD, BCORY, BCDRY, Home_Games)	-159.000000
126	(MOT, MOV, 01stD, D1stD, BCORY, BCDRY, Home_Ga...	-157.000000
108	(MOT, 01stD, D1stD, BCORY, BCDRY)	4.000000
79	(MOT, D1stD, BCORY, BCDRY)	4.717742
124	(MOT, 01stD, D1stD, BCORY, BCDRY, Home_Games)	4.864936
112	(MOT, D1stD, BCORY, BCDRY, Home_Games)	5.496399
73	(MOT, 01stD, D1stD, BCORY)	6.639582
109	(MOT, 01stD, D1stD, BCORY, Home_Games)	7.527562

37	(MOT, D1stD, BCORY)	9.070455
80	(MOT, D1stD, BCORY, Home_Games)	9.851029
74	(MOT, 01stD, D1stD, BCDRY)	10.209492
110	(MOT, 01stD, D1stD, BCDRY, Home_Games)	11.339412
33	(MOT, 01stD, D1stD)	11.902735
75	(MOT, 01stD, D1stD, Home_Games)	13.037000
38	(MOT, D1stD, BCDRY)	18.962838
81	(MOT, D1stD, BCDRY, Home_Games)	20.067068
9	(MOT, D1stD)	23.391556
39	(MOT, D1stD, Home_Games)	24.498137
111	(MOT, 01stD, BCORY, BCDRY, Home_Games)	41.730812
82	(MOT, BCORY, BCDRY, Home_Games)	42.453366
76	(MOT, 01stD, BCORY, BCDRY)	45.382767
40	(MOT, BCORY, BCDRY)	46.312955
78	(MOT, 01stD, BCDRY, Home_Games)	57.379042
35	(MOT, 01stD, BCDRY)	60.895260
77	(MOT, 01stD, BCORY, Home_Games)	66.867591
42	(MOT, BCDRY, Home_Games)	71.543445
41	(MOT, BCORY, Home_Games)	72.899644
34	(MOT, 01stD, BCORY)	72.986648
11	(MOT, BCDRY)	75.554218
10	(MOT, BCORY)	79.797014
36	(MOT, 01stD, Home_Games)	82.363182
93	(01stD, D1stD, BCORY, BCDRY)	84.791834
59	(D1stD, BCORY, BCDRY)	85.275023
118	(01stD, D1stD, BCORY, BCDRY, Home_Games)	85.553927
97	(D1stD, BCORY, BCDRY, Home_Games)	85.951332
8	(MOT, 01stD)	88.312222
53	(01stD, D1stD, BCORY)	95.640214
94	(01stD, D1stD, BCORY, Home_Games)	96.436377
22	(D1stD, BCORY)	99.009784
60	(D1stD, BCORY, Home_Games)	99.682168
54	(01stD, D1stD, BCDRY)	102.819514
95	(01stD, D1stD, BCDRY, Home_Games)	103.996779
12	(MOT, Home_Games)	109.775338
18	(01stD, D1stD)	111.883316
55	(01stD, D1stD, Home_Games)	113.070222
96	(01stD, BCORY, BCDRY, Home_Games)	115.822019
62	(BCORY, BCDRY, Home_Games)	116.307834
23	(D1stD, BCDRY)	117.161979
0	(MOT,)	117.289553
61	(D1stD, BCDRY, Home_Games)	118.309819
56	(01stD, BCORY, BCDRY)	119.165138
25	(BCORY, BCDRY)	119.844393
3	(D1stD,)	132.492463
24	(D1stD, Home_Games)	133.646960
58	(01stD, BCDRY, Home_Games)	146.097941

20	(01stD, BCDRY)	149.228845
57	(01stD, BCORY, Home_Games)	155.961367
19	(01stD, BCORY)	162.347453
26	(BCORY, Home_Games)	163.249544
27	(BCDRY, Home_Games)	166.600639
5	(BCDRY,)	170.282833
4	(BCORY,)	170.488632
21	(01stD, Home_Games)	187.701184
2	(01stD,)	193.870897
6	(Home_Games,)	229.700386

Based on the Mallows CP report we find that our best predictors of Margin of Victory (MOV) are Margin of Turnover (MOT) and Defensive 1st Downs Allowed (D1stD). We still observe very large Mallows CP values with these single explanatory variables relative to multivariable regression models.

6 3. Model Building: Explore several competing OLS models (based on part 2) and decide on one model only (with just one predictor). You will need to explain in detail how you arrived at your preferred model. Discuss the economic significance of your parameters, and overall findings. Make sure you discuss your main conclusions and recommendations. At a minimum, you need to include the following checks:

```
[35]: # Specify the Model usings results from Mallow cp
ols_mod = smf.ols(formula='MOV ~ MOT', data = eagles)

# Fit the Model
ols_fit = ols_mod.fit()

# Look at the Model Fit Summary
print(ols_fit.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          MOV      R-squared:                0.300
Model:                  OLS      Adj. R-squared:           0.296
Method:                 Least Squares      F-statistic:         73.21
Date:                  Tue, 24 Oct 2023      Prob (F-statistic):    6.43e-15
Time:                  23:15:28      Log-Likelihood:       -679.15
No. Observations:      173      AIC:                  1362.
Df Residuals:          171      BIC:                  1369.
Df Model:               1
Covariance Type:       nonrobust
=====
```


	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.7717	0.938	2.955	0.004	0.920	4.623
MOT	4.4975	0.526	8.556	0.000	3.460	5.535
=====						
Omnibus:		2.289	Durbin-Watson:			1.935
Prob(Omnibus):		0.318	Jarque-Bera (JB):			1.863
Skew:		0.222	Prob(JB):			0.394
Kurtosis:		3.247	Cond. No.			1.79
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

We fit a non robust model with an intercept of 2.77 (meaning we start off on average winning games with no turnovers by a margin of 3), followed by either a reduction or addition of 4 or 5 points for every turnover. This makes sense as turnovers can put offenses in advantageous positions, only having to attain a couple of first downs in order to achieve atleast 3 point, with the goal being 6 (but more often in practice 7).

```
[37]: # Evaluating the model
import statsmodels.stats.api as sms
from simple_colors import *

# Linearity: Harvey-Collier --> Ho: model is linear
name = ["t-stat", "p-value"]
test = sms.linear_harvey_collier(ols_fit)
print(blue("Linearity Test Results:", ['bold']))
print(list(zip(name, test)))
print("\n")

# Normaility of the Residuals: Jarque-Bera --> Residuals ~ N(0,1)
name = ["Jarque-Bera", "Chi^2 two-tail prob.", "Skew", "Kurtosis"]
test = sms.jarque_bera(ols_fit.resid)
print(blue("JB Results:", ['bold']))
print(list(zip(name, test)))
print("\n")

# Heteroskedasticity: Breush-Pagan --> Ho: var = constant
name = ["Lagrange multiplier statistic", "p-value", "f-value", "f p-value"]
test = sms.het_breuschpagan(ols_fit.resid, ols_fit.model.exog)
print(blue("BP Results:", ['bold']))
print(list(zip(name, test)))
```

Linearity Test Results:

[('t-stat', -0.9843242300354407), ('p-value', 0.3263632181550974)]

JB Results:

```
[('Jarque-Bera', 1.863410374413642), ('Chi^2 two-tail prob.',  
0.39388149572106407), ('Skew', 0.22229528943369903), ('Kurtosis',  
3.2466718767692595)]
```

BP Results:

```
[('Lagrange multiplier statistic', 0.3603567532061924), ('p-value',  
0.548308170650365), ('f-value', 0.3569342686266421), ('f p-value',  
0.5510041204231912)]
```

```
[38]: # Test for normality of residuals - JB test  
stats.jarque_bera(ols_fit.resid)
```

```
[38]: SignificanceResult(statistic=1.863410374413642, pvalue=0.39388149572106407)
```


Based on the output for the test for linearity, since the p-value is $> \alpha 0.05$ we fail to reject the null and conclude that the model is linear

As for normality test, the JB Results show that the errors are normally distributed because the p-value is $> \alpha 0.05$ even though the JB score = 1.8 which is > 1 (kurtosis is 3.2..., skewness is 0.2)

For Heteroskedasticity we observe that it is not possible to reject the null hypothesis that the variance is constant. This is due to the high p-value, and as such we can be fairly confident that as there is an increase in x there may be no observable increase in the variance of y .

```
[39]: # Diagnostic plots of ols fit  
figD = sm.graphics.plot_regress_exog(ols_fit, "MOT")  
figD.set_figheight(10)  
figD.set_figwidth(8)  
plt.show()
```

Regression Plots for MOT



MOV vs MOT shows a positive correlation. MOT vs residuals appears to be scattered around 0

```
[40]: model = smf.ols(formula='(MOV) ~ MOT + O1stD + BCORY + D1stD + BCDRY +_
      ↪Home_Games', data=eagles)
results = model.fit()
y = (eagles['MOV'])
y_pred=results.fittedvalues
```

```

# Using subset size =1
mr_sub = smf.ols(formula='(MOV) ~ MOT', data=eagles)
mr_sub_fit = mr_sub.fit()
y_sub=mr_sub_fit.fittedvalues

k = 8 # number of parameters in original model (includes y-intercept)
p = 3 # number of parameters in the subset model (includes y-intercept)

mallow.mallow(y, y_pred,y_sub, k, p)

```

[40]: 121.27263122025579

```

[41]: model = smf.ols(formula='(MOV) ~ MOT + O1stD + BCORY + D1stD + BCDRY +_
      ↪Home_Games', data=eagles)
results = model.fit()
y = (eagles['MOV'])
y_pred=results.fittedvalues

# Using subset size =1
mr_sub = smf.ols(formula='(MOV) ~ D1stD', data=eagles)
mr_sub_fit = mr_sub.fit()
y_sub=mr_sub_fit.fittedvalues

k = 8 # number of parameters in original model (includes y-intercept)
p = 3 # number of parameters in the subset model (includes y-intercept)

mallow.mallow(y, y_pred,y_sub, k, p)

```

[41]: 136.58084931793445

Using the results of the Mallows' CP the output showed that the best model with one predictor was MOV ~ MOT, followed by MOV~ D1stD. We tested the two models and found the first one had the lowest result, so we chose that as the best model with one predictor.

6.1 Look at Cook's distance Plot, Residuals Plot

```

[ ]: # Specify model
Turnover = smf.ols(formula='MOV ~ MOT', data=eagles)
T_results = Turnover.fit(cov_type='HC1')
print(T_results.summary())
T_results.resid.mean()

```


```

[43]: # Plot of residuals
sns.residplot(x='MOT', y='MOV', data=eagles,
              lowess=True, line_kws={'color':'red', 'lw':2, 'alpha':0.6})
plt.xlabel('Fitted Values')
plt.title('Residuals Plot')

```


```
plt.show
```

```
[43]: <function matplotlib.pyplot.show(close=None, block=None)>
```




We can find that towards the end of the graph we find significant high leverage observations in which the MOT shows a negative value but the MOV still remains quite high. These negative observations seem to heavily affect the residuals of the regression, even though it can be expected that a positive turnover margin would result a positive margin of victory.

```
[44]: # QQ plot on studentized residuals
student_resid=T_results.get_influence().resid_studentized
df=len(student_resid)-4
t_dist=stats.t(df)
sm.qqplot(student_resid, line='45', dist=t_dist)
plt.grid()
```




In this QQ plot of the studentized residuals we see 7 observations that fall outside an absolute standard deviation of 2 from the mean.

```
[45]: # Leverage plot
leverage=T_results.get_influence().hat_matrix_diag
plt.figure(figsize=(12,6))
plt.scatter(eagles.index,leverage)
plt.axhline(0,color='red')
plt.vlines(x=eagles.index, ymin=0,ymax=leverage)
plt.xlabel('Index')
plt.ylabel('Leverage')
plt.title('Diagnostic plot')
plt.show()
```




```
[46]: # Cook's distance plot
cooks_distance=T_results.get_influence().cooks_distance
plt.figure(figsize=(12,6))
plt.scatter(eagles.index,cooks_distance[0])
plt.axhline(0,color='red')
plt.vlines(x=eagles.index, ymin=0,ymax=cooks_distance[0])
plt.xlabel('Index')
plt.ylabel('Cook\'s Distance')
plt.title('Diagnostic plot')
plt.show()
```



The cook's distance graph shows one influential data point.

```
[47]: fig, ax= plt.subplots(figsize=(10,5))
fig=sm.graphics.influence_plot(T_results, ax=ax, criterion='cooks')
```



As we can see in the above influence and leverage plot there are a considerable amount of high leverage and high influence observations, but they can't be removed due to various reasons. These

include the fact that the games in which the Eagles play are both real data points and that there are deliberate decisions made by professionals that go into the outcomes of football games. As such it would not be appropriate to prune the data of high leverage data.

6.2 Evaluate the robustness of your coefficient estimates by bootstrapping your model. Provide a histogram of the bootstrapped estimates (including R^2), and comment on the findings. In particular how do these estimates compare against your LS estimates?


```
[48]: from sklearn.linear_model import LinearRegression
```

```
[49]: coefs=pd.DataFrame(columns=['B0','B1'])

for i in range(1000):
    sample = eagles.sample(eagles.shape[0], replace=True)
    results=smf.ols('MOV~MOT', sample).fit()
    b0,b1=results.params
    coefs=coefs._append({'B0':b0, 'B1':b1}, ignore_index=True)
# 5% CI
b0_u, b1_u = coefs.quantile(.975)
b0_l, b1_l = coefs.quantile(.025)
```


```
[50]: # Bootstrapped Histograms with CI
coefs.B0.hist()
plt.xlabel("Bootstrap Intercepts")
plt.ylabel("Frequency")
plt.axvline(b0_u, color = "red")
plt.axvline(b0_l, color = "red")
```

```
[50]: <matplotlib.lines.Line2D at 0x233cfed2b90>
```



```
[51]: coefs.B1.hist()  
plt.xlabel("Bootstrap B1")  
plt.ylabel('Frequency')  
plt.axvline(b1_u, color = 'red')  
plt.axvline(b1_l, color = 'red')
```

```
[51]: <matplotlib.lines.Line2D at 0x233cfae4210>
```



Based on the bootstrap results the intercepts (about 2.8) and the slopes (about 4.5) are close enough to those from the results of the LS estimates (2.77 and 4.49)

6.3 Confidence intervals

```
[52]: from scipy.stats import bootstrap

# Defining MOT coefficient
def reg_boot_b1(x,y):

    x = x.reshape((len(x),1))
    y = y.reshape((len(x),1))
    reg = LinearRegression().fit(x,y)

    return reg.coef_[0][0]
    #Pulling out Beta1 value

# Defining intercept
def reg_boot_intercept(x,y):

    x = x.reshape((len(x),1))
```

```

y = y.reshape((len(x),1))
reg = LinearRegression().fit(x,y)

return reg.intercept_[0]
# Pulling out Intercept value

X = eagles.MOT
Y = eagles.MOV
res = bootstrap((X,Y), reg_boot_b1, confidence_level=0.95, vectorized=False,
↳method='BCa',
        paired=True)
print(res.confidence_interval)

```

ConfidenceInterval(low=3.5597200129169444, high=5.5465785672423324)

With a 95% CI we are confident that the estimated reg_boot_b1 (beta 1) falls between (3.540241015898647, 5.526534689378149)

```

[53]: res = bootstrap((X,Y), reg_boot_intercept, confidence_level=0.95,
↳vectorized=False, method='BCa',
        paired=True)
print(res.confidence_interval)

```

ConfidenceInterval(low=0.9265648020958546, high=4.6401653396316025)

With a 95% CI we are confident that the estimated reg_boot_intercept (beta 0) falls between (1.0055391051206657, 4.604349878033104)

6.4 Use cross-validation to evaluate your model's performance

```

[54]: import numpy as np
from sklearn.model_selection import train_test_split
from sklearn import linear_model
from sklearn.linear_model import LinearRegression
from sklearn import metrics
from sklearn.model_selection import cross_val_score

x = eagles[['MOT']]
y = eagles[['MOV']]
# Perform an OLS fit using all the data
regr = LinearRegression()
model = regr.fit(x,y)
regr.coef_
regr.intercept_

# Split the data into train (70%)/test(30%) samples:
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3,
↳random_state=0)

```

```

# Train the model:
regr = LinearRegression()
regr.fit(x_train, y_train)

# Make predictions based on the test sample
y_pred = regr.predict(x_test)

# Evaluate Performance
print('MAE:', metrics.mean_absolute_error(y_test, y_pred))
print('MSE:', metrics.mean_squared_error(y_test, y_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))

# Perform a 5-fold CV
regr = linear_model.LinearRegression()
scores = cross_val_score(regr, x, y, cv=5,
    ↳scoring='neg_root_mean_squared_error')
print('5-Fold CV RMSE Scores:', scores)

```

```

MAE: 9.435276975984142
MSE: 162.93430541061392
RMSE: 12.764572276837713
5-Fold CV RMSE Scores: [-12.89987682 -10.67071177 -12.22473686 -11.72389117
-14.56047974]

```

Root Mean Squared Error (RMSE), which measures the average prediction error made by the model in predicting the outcome for an observation shows a value of 12.765. For the variable MOT, the mean is 2.9, the median is 3, and Std is 14.

The lower the RMSE, the better the model. In this case we have a high RMSE which may mean that our model does not accurately explain all the dynamics that result in values we see in the MOV column. Our R^2 of 0.300 with an adjusted R^2 of 0.296 shows that initially regressing MOV on MOT only accounted for about 30% of the observed dynamics between the two variables.

Ultimately, we see that our value of RMSE falls just under the standard deviation and as a result we could be anywhere from 10-13 points off from the actual MOV. That is a value of 1.5 touchdowns to almost 2 touchdowns, which is quite large in terms of football as teams scores an average of 3.5 touchdowns per game in the modern NFL. As such it shows that we need more than one explanatory variable when it comes to truly analyzing MOV.

With that being said, ball possession and security are considered tantamount to success in the NFL. When the Eagles went on their 8-game winning streak in 2022, many highlighted the large positive turnover margin as a key to their success. But one could also see that the Eagles offense was excellent in terms of Red Zone Offense and large passing plays. These dynamics can't be shown, and also, it's questionable whether they would be necessary in a proper regression analysis due to their relation with the MOV variable and the ensuing endogeneity.

Overall, MOT is statistically significant when it comes to determining game outcomes in the form of MOV. But, using one variable is simply not enough to be able to predict something so complex and variable, especially when there are aspects of “relative” luck, skill, psychology, time of games, offensive prowess etc. MOT is a variable that we can use as a step to forming a more accurate

multivariable regression, but can't explain everything that gives the resulting MOV.

So, while the RMSE values in comparison to our data may make it seem that using MOT to predict MOV may seem dubious, the real issue comes from the fact that the model itself is not robust due to the lack of predictors. MOT is both valuable as a way to predict MOV and overall team success in both layman's and a statistician's terms, but it is just one of many confounding variables that determine on field performance in the game of football. You keep turnovers down on the offensive side of the ball, and get turnovers on the defensive side of the ball the more likely it is you are going to score and beat your opponent. The more you do this the greater the expected MOV, but this is not always the case, and that makes America's sport so enigmatic in nature.